

# K HARI KRISHNA

COMPUTER VISION ENGINEER

+91 9177595499

smarth326@gmail.com

Mahbubnagar, Telangana, India



## SUMMARY

Efficient and hands-on AI Systems Engineer with strong expertise in **AI model deployment**, **TensorRT-based optimization**, and **image/video processing** for real-time applications. Skilled in integrating pretrained deep learning models, **accelerating inference pipelines**, and **reducing system latency through lightweight engineering solutions**. Experienced in transforming traditional pipelines into faster, resource-efficient systems using tools like Docker, Pillow, and multiprocessing, even without deep customization of deep learning frameworks.

## EDUCATION

### Bachelor of Commerce (B.Com) in Computer Applications

Government Degree College, Khairatabad, India  
2019-2022

## TECHNICAL SKILLS

- Languages:** Python (Strong), Basic C++.
- Model Integration:** TensorRT, YOLOv7 (Pretrained)
- Libraries/Tools:** OpenCV, Pillow, FFmpeg, MediaMTX, PyTorch (basic usage).
- Systems & Tools:** Docker, RabbitMQ, Redis, Git, Label Studio, Django (API-level)
- Performance Optimization:** Parallel Processing (Multithreading, Multiprocessing), Real-Time Pipeline Efficiency

## ACHIEVEMENTS

- Awarded "**Innovator of the Month**" for successfully optimizing platform components and contributing impact R&D solutions.
- Played a key role in bench-marking and transitioning major workloads to GPU.

## PROFESSIONAL EXPERIENCE

### Computer Vision Intern → Computer Vision Engineer

ZestloT Technologies Pvt. Ltd., Hyderabad, Telangana, India.  
Feb 2023 – Present

- Started as a **Computer Vision Intern** in Feb 2023 and converted to a full-time **Computer Vision Engineer** in Sep 2023 after consistently contributing to high-impact R&D and production solutions.
- Integrated **YOLOv7-based object detection** and **segmentation** models into the company's AI platform.
- Performed **TensorRT optimization** for object detection and segmentation models, significantly improving inference speed and reducing hardware utilization.
- Converted pretrained models from **PyTorch → ONNX → TensorRT** and deployed them in real-time systems.
- Replaced OpenCV with Pillow to accelerate image processing, achieving a **2x performance boost** and lower memory usage.
- Built and deployed an **image augmentation pipeline** as a **Dockerized microservice** integrated with the platform.
- Resolved threading and multiprocessing issues to enable efficient, parallel execution in real-time pipelines.
- Resolved library and package conflicts to successfully** run detection and segmentation models in a single production container, reducing memory usage, startup time, and deployment complexity.